

**CS 521:** Systems Programming

# Proof-of-Work and Bitcoin

Lecture 17

# Today's Agenda

---

- Proof-of-work Systems
- Hashcash
- Bitcoin

# Today's Agenda

---

- **Proof-of-work Systems**
- Hashcash
- Bitcoin

# Proof-of-Work [1/2]

- *Wait a minute! I thought we were learning about parallel programming!*
  - Did Matthew forget what class this is?
- Well, friends... Project 3 is up!
  - And it involves threads
  - And bitcoin mining

# Proof-of-Work [2/2]

---

- Proof-of-work (POW) systems help prevent DDoS attacks and other types of spamming
- Also useful in cryptocurrencies
- Main idea: give up some of your time (or computational power) to **legitimize** an action or object

# Shell Money

---

- Sea shells were used for thousands of years as legal tender
- It takes time to collect shells, carve them, etc.
  - In some cases, the shells were woven into fabric/leather
  - The currency itself reflected the time it took to be made, and therefore determined its value
- Different groups used different shells/designs
  - Only carry value because we say so

# CAPTCHA

- **C**ompletely **A**utomated **P**ublic **T**uring test to tell **C**omputers and **H**umans **A**part
- CAPTCHAs are basically proof-of-work systems for humans
- So in other words, POW is an annoying, time consuming task for your computer to do just in the interest of proving it's not spamming/DDoSing
  - Luckily computers don't get annoyed as easily as we do...

# Pricing Functions

- POW systems use *pricing functions* to give the computer a workout
  - **Dwork C., Naor M.** *Pricing via Processing or Combatting Junk Mail.*
- A pricing function  $f$  has the following requirements:
  - $f$  is moderately easy to compute
  - $f$  is not amenable to amortization: given  $L$  values,  $m_1 \dots m_L$  the amortized cost of computing  $f(m_1) \dots f(m_L)$  is comparable to computing  $f(m_i)$  for any  $1 \leq i \leq L$
  - given  $x$  and  $y$ , it is easy to determine if  $y = f(x)$



# Hash Inversions

- A common pricing function is having the computer perform **hash inversions**
  - “What was the **input** that produced this hash code?”
- Hash inversions are tough to compute (assuming a cryptographic hash function)
  - After all, they’re designed to be one way functions
  - Any time we map an infinite set of inputs to a finite set of numbers (hash space), this is feasible, but still tough.

# An Example [1/2]

- Let's say our mission is to find a hash with four leading zeros
- Start out with what we want to send:
  - "Hello World!"
- We also need to append a **nonce**
  - *Number used only once*
  - We increase this with each hash attempt
  - This will change our output hash each iteration

# An Example [2/2]

- This approach allows us to eventually find our matching hash, but has a weakness
  - We can precompute the hashes and re-use them later
- We also need some type of identifier for this particular transaction
  - Maybe a centralized service hands out transaction IDs
  - We could use the current time, as long as we can assume clocks are reasonably synced up

# A Pricing Function

```
while True:
    nonce = nonce + 1
    string = message + str(nonce)
    hash = sha1(string)
    if prefix(hash) == '0000':
        # Send message with hash
        break
```

# Verification

---

```
if sha1(msg.payload) == msg.hash:  
    # It's valid... Whew! That was tough!  
    # (You could also verify the  
    # transaction id or timestamp here)
```

# Varying the Difficulty

- To change the difficulty, we'll just adjust the number of zeros we want
- Unfortunately, the difficulty won't increase linearly
- Approaches:
  - Perform a bitwise comparison rather than string (allows more precision)
  - Have the sender perform multiple inversions(maybe message1 + another nonce)

# Proof-of-Work Variants

---

1. Challenge-response
2. Solution verification

# Challenge-Response

---

- Interactive link between sender and receiver
- Receiver chooses the challenge
  - Can adapt the challenge based on its own load or remote capabilities
- After completing the challenge, the service can be accessed by the sender



# Solution Verification

---

- In this case, we already know what the challenge is in advance
  - Compute it locally and send the message
- Receiver needs to verify the message and then process it
- Better for one-time communication (like email)

# I/O Challenge

---

- Hardware improvements in this space (disks, memory, network, etc) have traditionally been less rapid
  - Helps us avoid reconfiguring the difficulty frequently
- This could be a memory-intensive algorithm
- Another approach is requesting tokens from remote servers
  - Here the work is "tough" because of latencies rather than the computation

# Today's Agenda

---

- Proof-of-work Systems
- **Hashcash**
- Bitcoin

# Hashcash

- Adds a new header to emails in an effort to reduce DDoS/spam
  - Proposed by Dwork and Naor
- You can install the hashcash command line utility:

```
$ hashcash -m 'mmalensek@usfca.edu'
```

```
hashcash stamp: 1:20:200927:mmalensek@usfca.edu::ZeNi
```

```
+DkIeFrH3aU1:000000000000000000000000000000000000000000000
```

009nf0

- On the receiving end, all that has to be done is verify the SHA-1 hash of the header

# Header Fields

[illegible]

- ver: Hashcash version
- bits: Number of zero bits
- date: The time that the message was sent: YYMMDD
- resource: Resource data string being transmitted
- ext: Extension (currently ignored)
- rand: String of random characters
- counter: Nonce

# Sending a Message

---

- The sender performs the hash inversion and prepares the header
  - This takes a little CPU time, but shouldn't be noticeable
- Adds the header to the email message
- Performs the send operation as usual

# Receiving a Message

---

- On the receiving side, all we need to do is compute the SHA-1 hash of the entire Hashcash header
- Then we check:
  - That the correct number of leading zeroes is present
  - The provided date is valid
- This takes an imperceptible amount of time

# Why Hashcash Works

---

- Even heavy email users only send a few hundred emails per day
- Spammers want to send millions
  - This is going to cost a lot of CPU time
- Additionally, sending an email with no header or an incorrect header will incur steep penalties
  - Too many incorrect headers? Ban the IP
- Best of all, we don't have to start paying for email



# Why it doesn't work [1/2]

- Back in 1992 when Hashcash was invented, we didn't have such a huge variety of computing hardware
  - Smartphones, tablets, refrigerators, etc.
  - This makes coming up with the right difficulty for the challenge... difficult.
- The power of computing hardware isn't distributed uniformly across the Earth
- Hash inversions are amenable to parallelism and custom hardware

# Why it doesn't work [1/2]

- Spammers could adopt similar hardware to that of Bitcoin miners
  - GPUs, ASICs
  - Depends on cost vs. benefit
  - Related: cloud instances. Computing is so cheap!
- Since email is decentralized, you can't force everyone to use this new standard
  - Would actually be easier nowadays (get Google and Microsoft on board, and you're just about done)

# Today's Agenda

---

- Proof-of-work Systems
- Hashcash
- **Bitcoin**

# Bitcoin Now

- 1 BTC = 28,667 USD 🤖
  - Every time I talk about bitcoin I realize how rich I would've been if I'd have bought some in 2017
- ~600,000 transactions per day
  - Not a massive increase over the years
- ~19m bitcoins in circulation
  - Maxes out at 21m
- See: <http://blockchain.info>

# Blockchain

---

- The Bitcoin **blockchain** is a decentralized database of Bitcoin transactions
- Each block in the chain includes the hash of the previous block
- Starts with the **genesis block**
- When a transaction occurs, it is added to the current block and will be verified by miners

# Blocks

- A **block** is a list of transactions with some metadata
- Magic number (4 bytes) = 0xD9B4BEF9
- Block size (4 bytes)
- Block header
- Transaction counter
- Transaction data

# Block Headers

---

- Version
- Hash of the previous block
  - This makes tampering with the chain difficult
- Current hash of the transactions in the block
- Timestamp (last update)
- Difficulty
- Nonce

# Reaching Agreement

---

- Bitcoin allows forks off of the current block
- Whichever fork is acknowledged and used by the most participants becomes the “true” path
  - Longest path wins
- Transactions that went to a “failed” fork are added back to the “true” blockchain



# Reaching an Agreement

---

- Provisions are in place to ensure transactions are dealt with in a reasonable amount of time
  - Target: 10 minutes
- Every 2,016 blocks the system automatically adjusts its difficulty to hit the 10-minute target
- From 2014 – 2015 the average number of nonces tried before a new block could be created increased from 16 quintillion to 200 quintillion

# Mining Bitcoin

- Bitcoin uses the Hashcash algorithm for a different purpose: mining coins
- "Mining" means verifying a block of transactions
  - Finding the nonce (aka **solution**)
- Miners, who are the basis of transaction verification, are paid in new bitcoins and transaction fees
  - The reward of new bitcoins is halved every 210,000 blocks (~4 years)
  - Monetary supply limited to 21m bitcoins

# Verification

- In bitcoin, the **difficulty** of the challenge is varied to keep the network chugging along
- Once all 21m bitcoins are created, miners will be rewarded for verification via transaction fees only
- What is the cost vs. benefit of mining these coins?
  - Electricity vs. the size of the reward
- Lots of companies now build power-efficient hardware specifically for mining

# Pooled Mining

- As difficulty goes up, the chances of a single miner verifying a block goes down
- To combat this, **pools** of miners formed
- Pools divide up the work (nonces) among participants
  - Rewarded with a share of new bitcoins based on how much work was done
  - Less wasted effort, but less reward

# Moral Issues

---

- We are consuming massive amounts of fossil fuels to produce fake money
  - Production is only hard because we make it so
- Mining hardware gets bought up and then discarded once we move to harder hash inversions
- Some Useful Proof-of-Work systems try to do beneficial work
  - Finding prime numbers (Primecoin)
  - Protein folding (Curecoin)